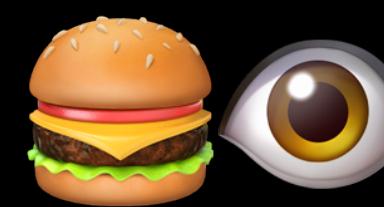




TensorFlow

Milestone Project 1



Food Vision Big™

Where can you get help?

- Follow along with the code



```
In the previous notebook (transfer learning part 3: scaling up) we built Food Vision mini: a transfer learning model which beat the original results of the Food101 paper with only 10% of the data.  
But you might be wondering, what would happen if we used all the data?  
Well, that's what we're going to find out in this notebook!  
We're going to be building Food Vision Big™, using all of the data from the Food101 dataset.  
Yep. All 75,750 training images and 25,250 testing images.  
And guess what...  
This time we've got the goal of beating DeepFood, a 2016 paper which used a Convolutional Neural Network trained for 2-3 days to achieve 77.4% top-1 accuracy.  
  
Note: Top-1 accuracy means "accuracy for the top softmax activation value output by the model"  
(because softmax outputs a value for every class, but top-1 means only the highest one is evaluated).
```

"If in doubt, run the code"

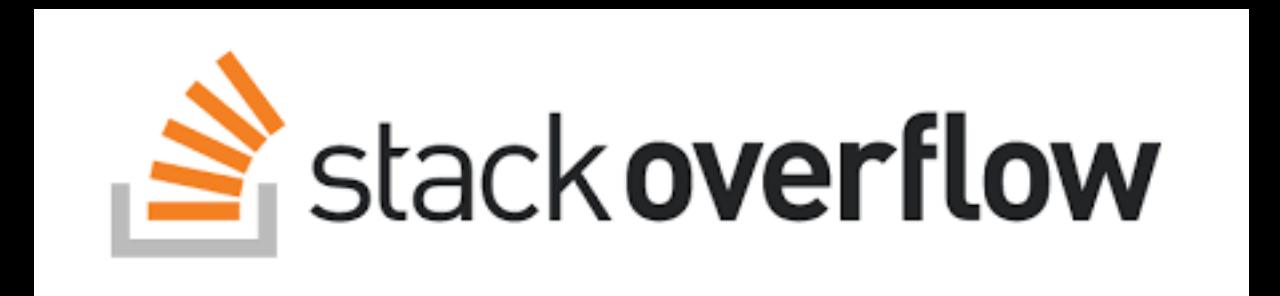
- Try it for yourself



- Press SHIFT + CMD + SPACE to read the docstring

```
def load(name: str, *, split: Optional[Tree[splite.Lib.Split]] = None, data_dir: Optional[str] = None, batch_size: Optional[int] = None, shuffle_files: bool = False, download: bool = True, as_supervised: bool = False, decoders: Optional[TreeDict[decode.Decoder]] = None, read_config: Optional[read_config.Lib.ReadConfig] = None, with_info: bool = False, builder_kwargs: Optional[Dict[str, Any]] = None, download_and_prepare_kwargs: Optional[Dict[str, Any]] = None, as_dataset_kwargs: Optional[Dict[str, Any]] = None, try_gcs: bool = False) -> Dataset: Loads the named dataset into a tf.data.Dataset.  
tfds.load is a convenience method that:  
  
(train_data, test_data), ds_info = tfds.load(name="food101", # target dataset to get from TFD  
split=[["train", "validation"], # what splits of  
shuffle_files=True, # shuffle files on download?  
as_supervised=True, # download data in tuple for  
with_info=True) # include dataset metadata if s  
  
Downloading and preparing dataset food101/2.0.0 (download: 4.65 GiB, generated: Unknown size,  
DI Completed... 100% 1/ [06:05<00:00, 365.33s/ url]  
DI Size... 100% 4764/4764 [06:05<00:00, 13.04 MiB/s]  
Extraction completed... 100% 1/ [06:05<00:00, 365.26s/ file]
```

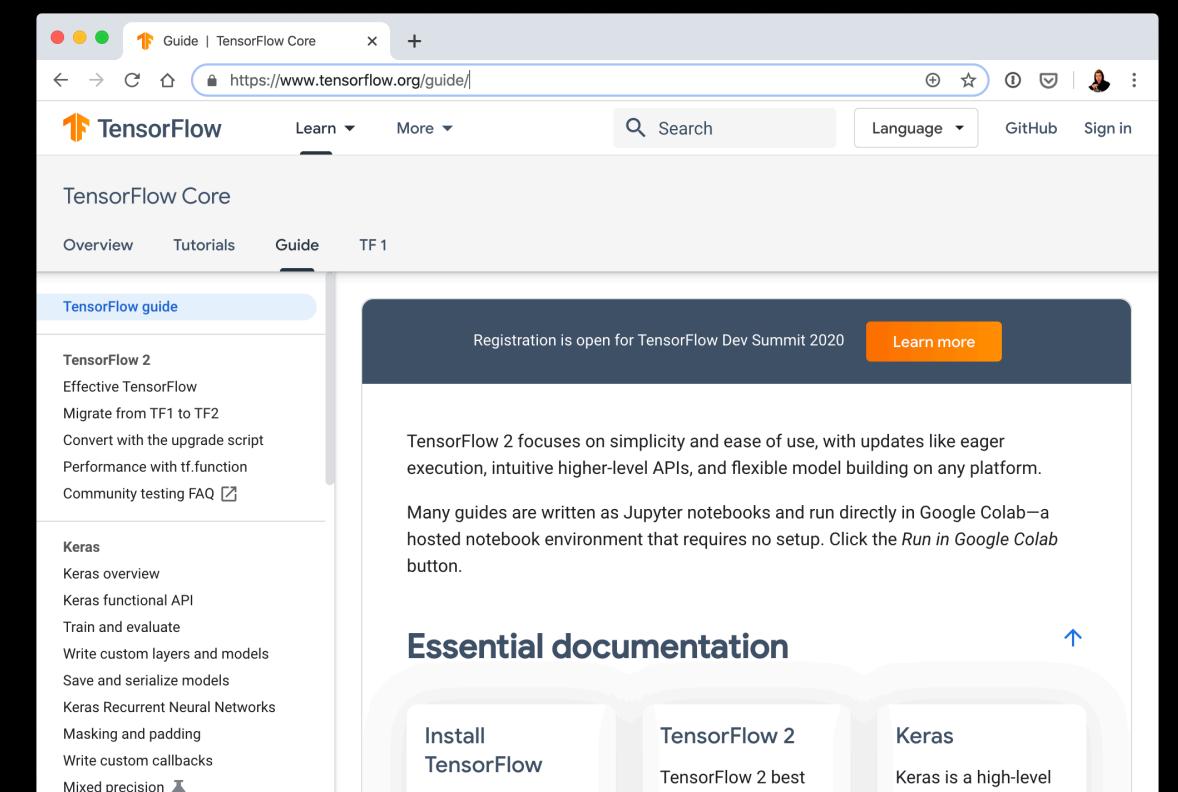
- Search for it



- Try again

- Ask (don't forget the Discord chat!)

(yes, including the "dumb" questions)



A new way of approaching things

What we're going to cover

(broadly)

- Using TensorFlow Datasets to download and explore data (all of Food101)
- Creating a **preprocessing function for our data**
- Batching & preparing datasets for modelling (**making them run fast**)
- Setting up **mixed precision training** (**faster model training**)

What you're going to cover...

- Building and training a feature extraction model
- Fine-tuning your feature extraction model to the beat the DeepFood paper 📈
- Evaluating your model results on TensorBoard
- Evaluating your model results by **making and plotting predictions**

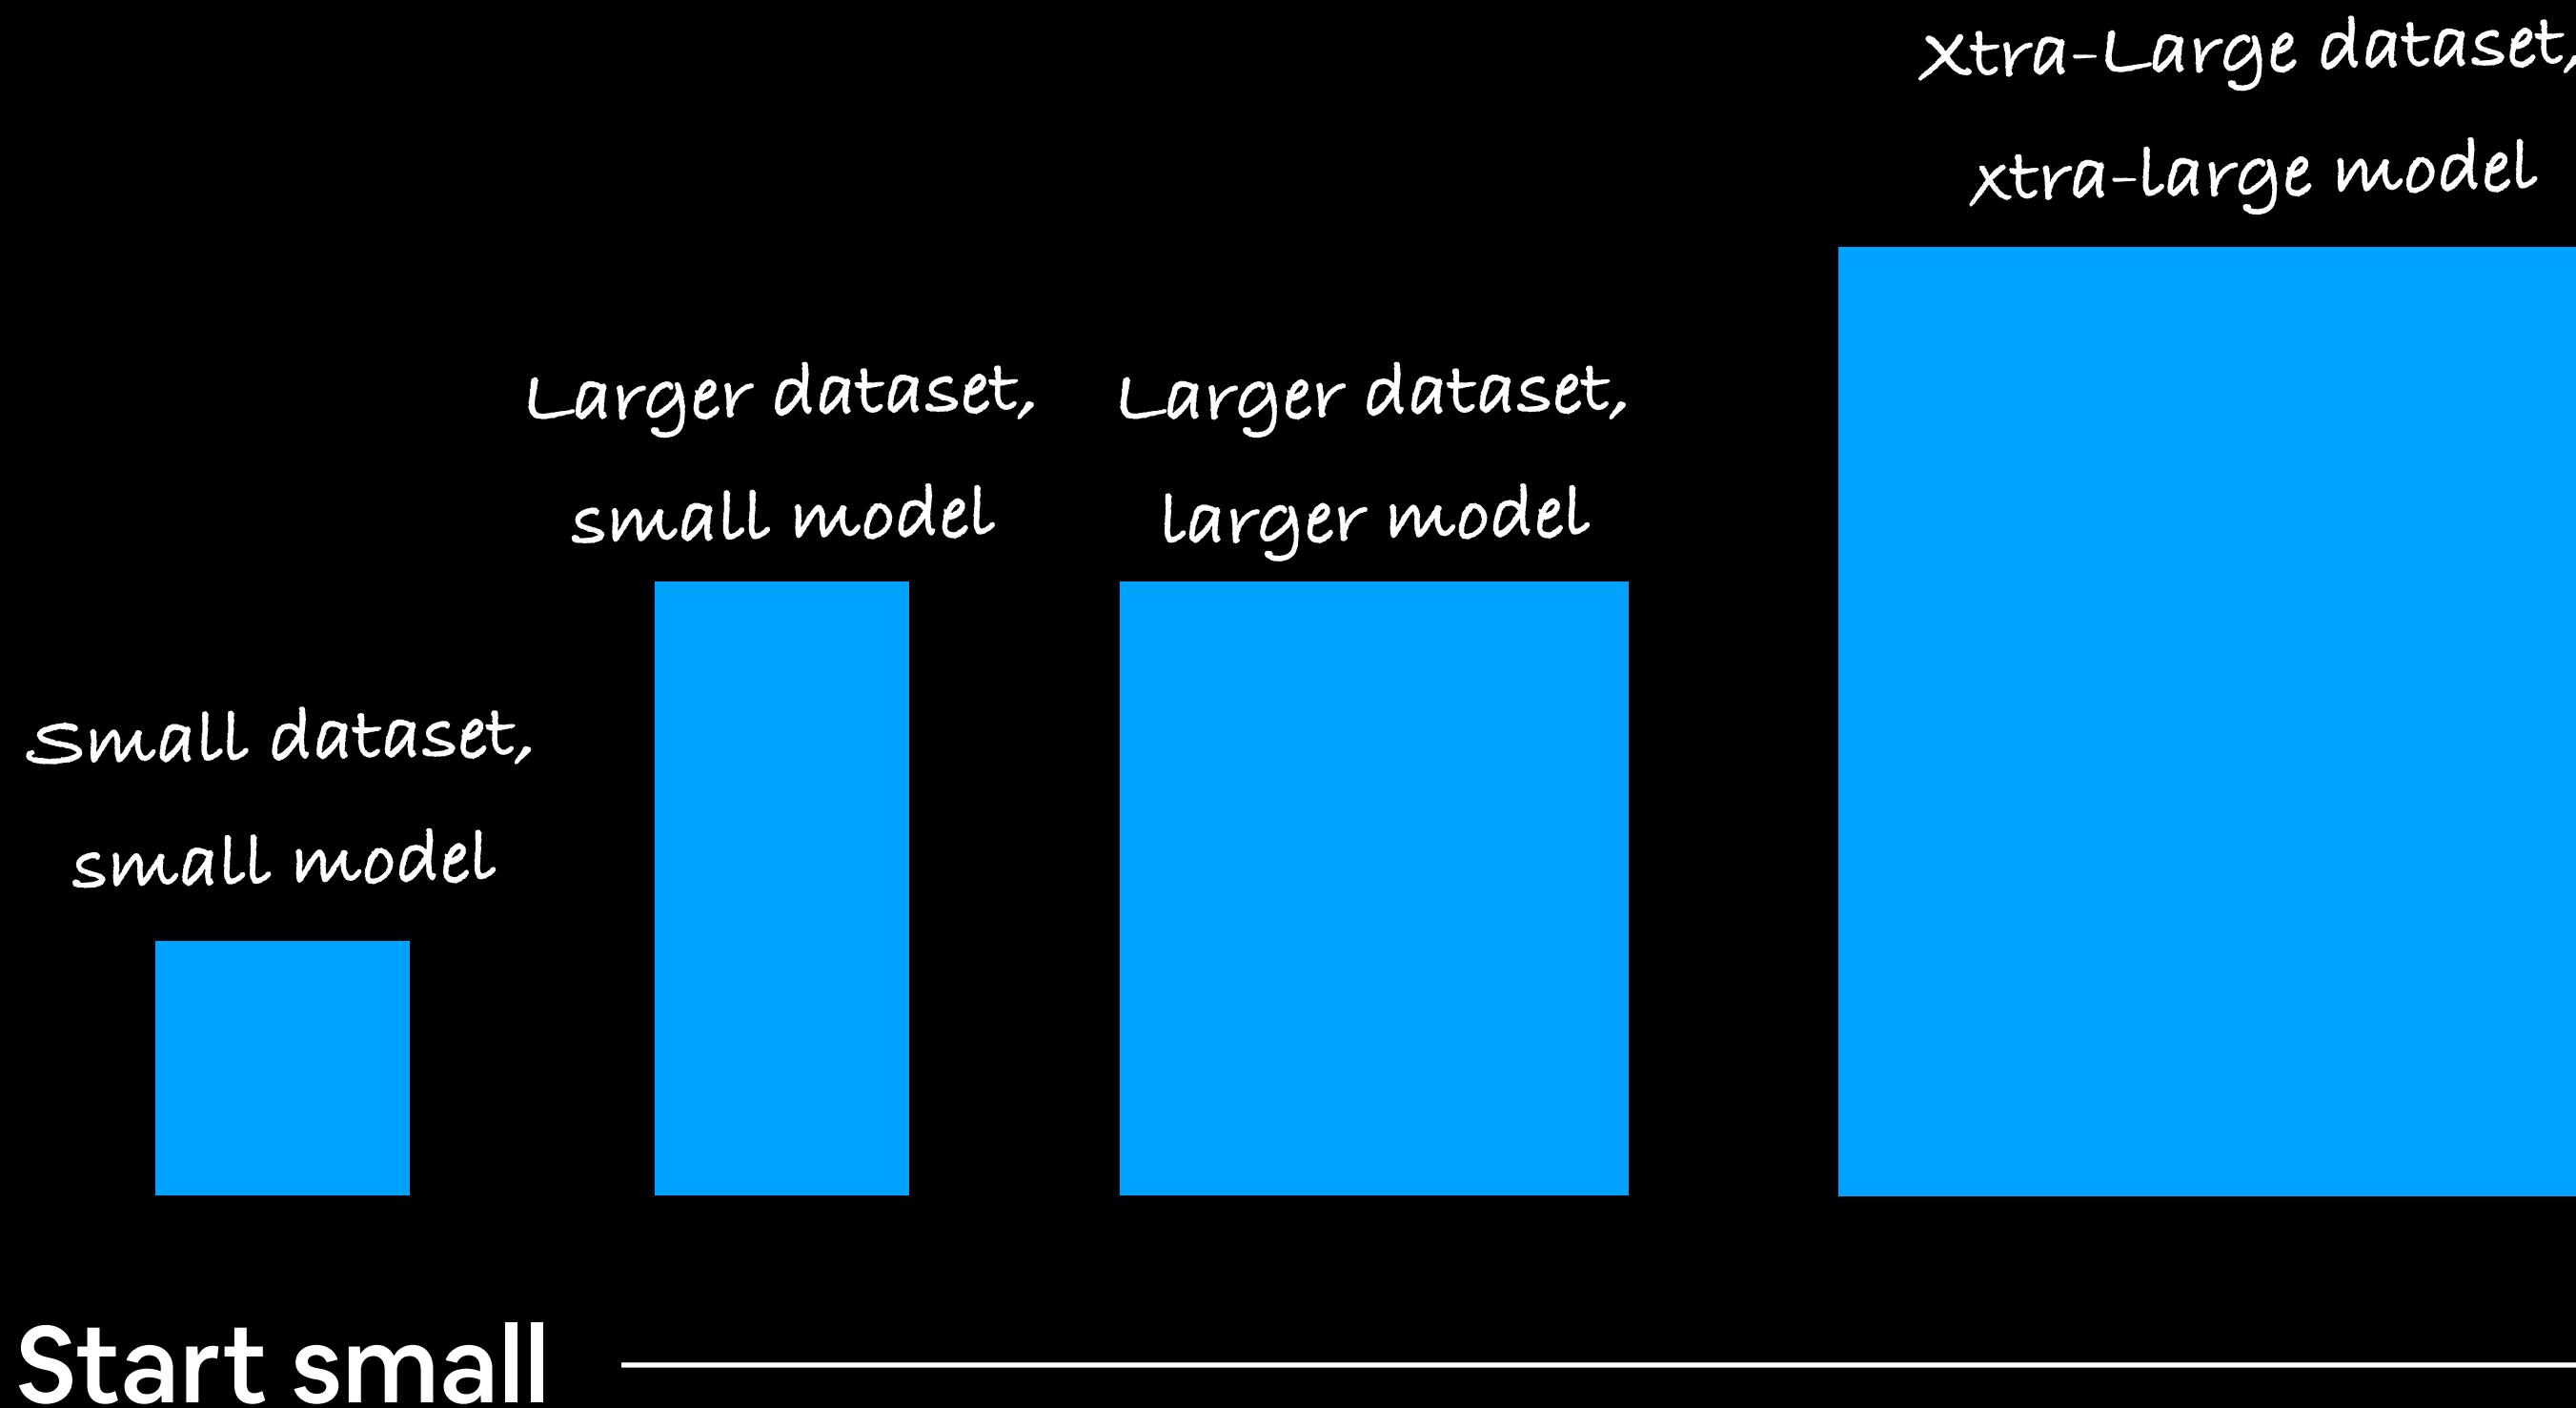
(cook up lots of code!)

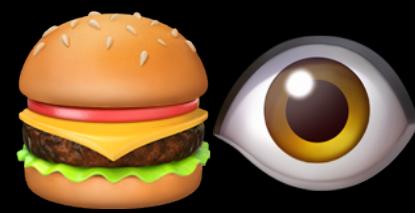
How:



Serial experimentation

Everything you got





Food Vision: Dataset(s) we're using

Note: For randomly selected data, the Food101 dataset was downloaded and modified using the [Image Data Modification Notebook](#)

Dataset Name	Source	Classes	Training data	Testing data
pizza_steak	Food101	Pizza, steak (2)	750 images of pizza and steak (same as original Food101 dataset)	250 images of pizza and steak (same as original Food101 dataset)
10_food_classes_1_percent	Same as above	Chicken curry, chicken wings, fried rice, grilled salmon, hamburger, ice cream, pizza, ramen, steak, sushi (10)	7 randomly selected images of each class (1% of original training data)	250 images of each class (same as original Food101 dataset)
10_food_classes_10_percent	Same as above	Same as above	75 randomly selected images of each class (10% of original training data)	Same as above
10_food_classes_100_percent	Same as above	Same as above	750 images of each class (100% of original training data)	Same as above
101_food_classes_10_percent	Same as above	All classes from Food101 (101)	75 images of each class (10% of the original Food101 training dataset)	250 images of each class (same as original Food101 dataset)
“food101”	Same as above (from TensorFlow Datasets)	All classes from Food101 (101)	75,750 images (750 images per class)	250 images of each class (same as original Food101 dataset)

Let's code!

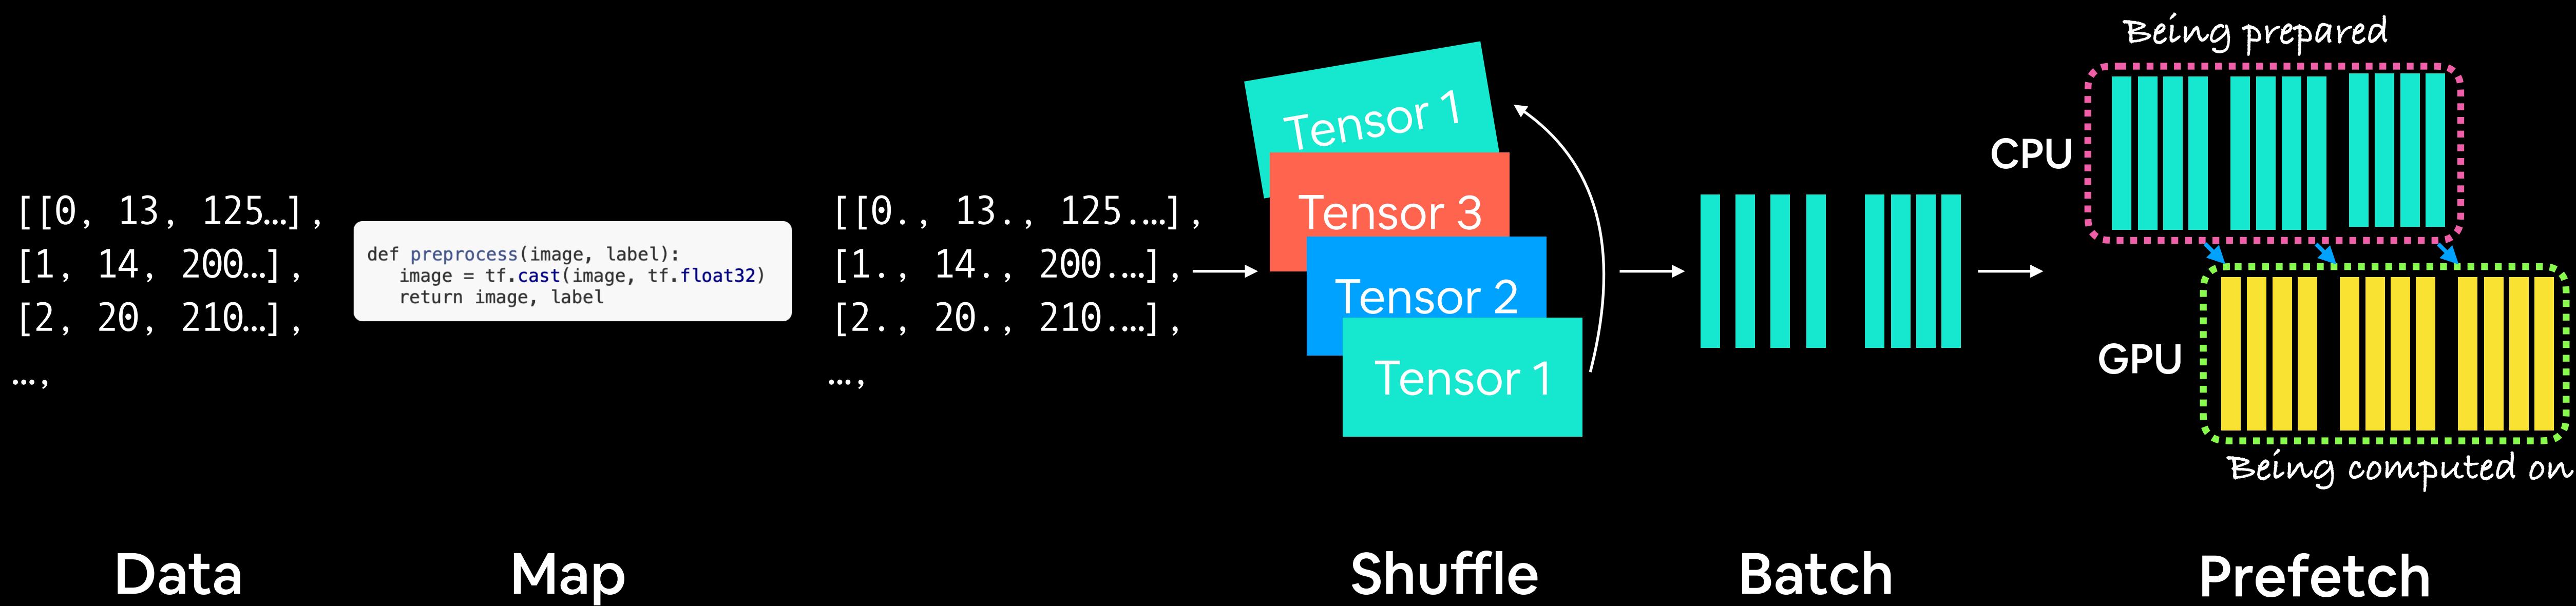
What is TensorFlow Datasets (TFDS)?

TensorFlow Datasets is a place for prepared and ready-to-use machine learning datasets.

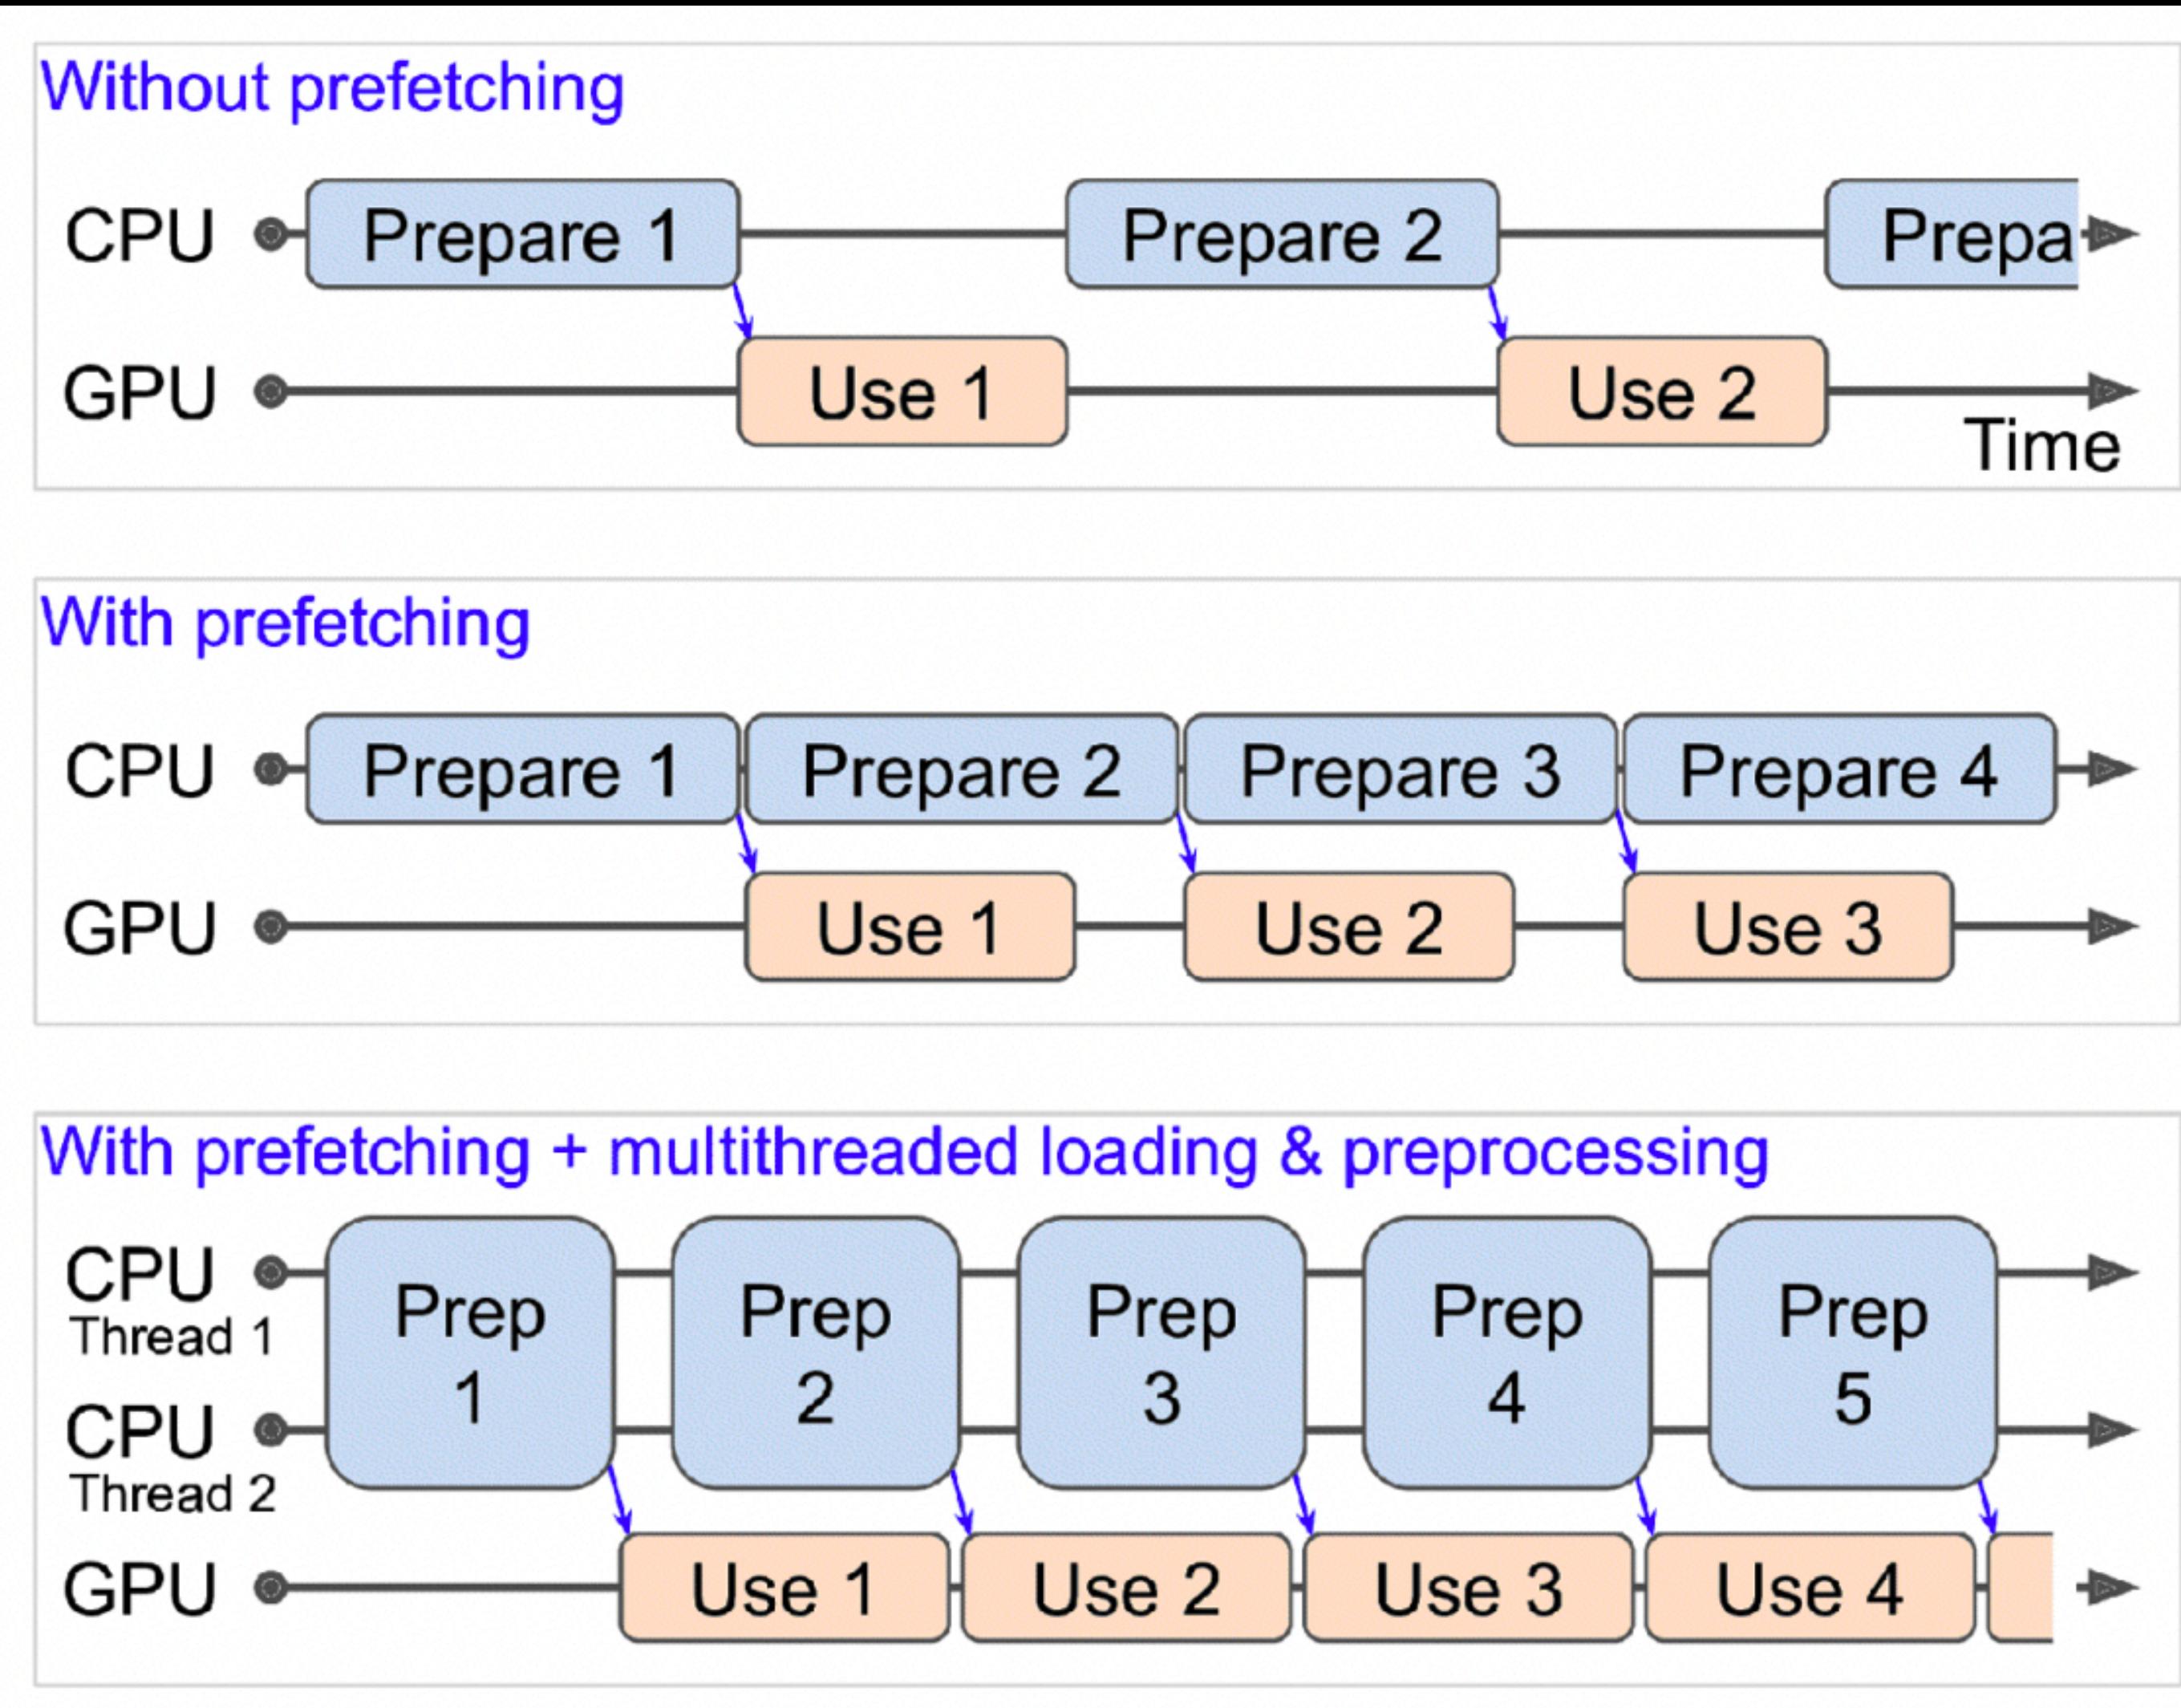
- Why use TensorFlow Datasets?
 - Load data already in tensor format
 - Practice on well established datasets (for many different problem types)
 - Experiment with different modelling techniques on a consistent dataset
 - Why **not** use TensorFlow Datasets?
 - Datasets are static (do not change like real-world datasets)

Batching & preparing datasets

```
# Map preprocessing function to training data (and parallelize)
train_data = train_data.map(map_func=preprocess, num_parallel_calls=tf.data.AUTOTUNE)
# Shuffle train_data and turn it into batches and prefetch it (load it faster)
train_data = train_data.shuffle(buffer_size=1000).batch(batch_size=32).prefetch(buffer_size=tf.data.AUTOTUNE)
```



Prefetching



Source: Page 422 of [Hands-On Machine Learning with Scikit-Learn, Keras & TensorFlow Book](#) by Aurélien Géron

What are callbacks?

- Callbacks are a tool which can **add helpful functionality** to your models during training, evaluation or inference
- Some popular callbacks include:

Callback name	Use case	Code
<u>TensorBoard</u>	Log the performance of multiple models and then view and compare these models in a visual way on TensorBoard (a dashboard for inspecting neural network parameters). Helpful to compare the results of different models on your data.	<code>tf.keras.callbacks.TensorBoard()</code>
<u>Model checkpointing</u>	Save your model as it trains so you can stop training if needed and come back to continue off where you left. Helpful if training takes a long time and can't be done in one sitting.	<code>tf.keras.callbacks.ModelCheckpoint()</code>
<u>Early stopping</u>	Leave your model training for an arbitrary amount of time and have it stop training automatically when it ceases to improve. Helpful when you've got a large dataset and don't know how long training will take.	<code>tf.keras.callbacks.EarlyStopping()</code>

What is mixed precision training?

Mixed precision training uses the combination of float32 and float16 data types to speed up model training (can result in a 3x speed up on modern GPUs).



No mixed precision



With mixed precision*

*3x speedup = maybe

Before and after mixed precision

```
import tensorflow as tf
from tensorflow.keras import layers

# Download base model and freeze underlying layers
base_model = tf.keras.applications.EfficientNetB0(include_top=False)
base_model.trainable = False

# Create functional model
inputs = layers.Input(shape=INPUT_SHAPE)
x = base_model(inputs, training=False)
x = layers.GlobalAveragePooling2D()(x)
outputs = layers.Dense(10, activation="softmax")(x)
model = tf.keras.Model(inputs, outputs)

# Compile
model.compile(loss="sparse_categorical_crossentropy",
              optimizer=tf.keras.optimizers.Adam(),
              metrics=["accuracy"])
```

No mixed precision

```
import tensorflow as tf
from tensorflow.keras import layers
from tensorflow.keras import mixed_precision 1

# Turn on mixed precision training
mixed_precision.set_global_policy("mixed_float16") 2

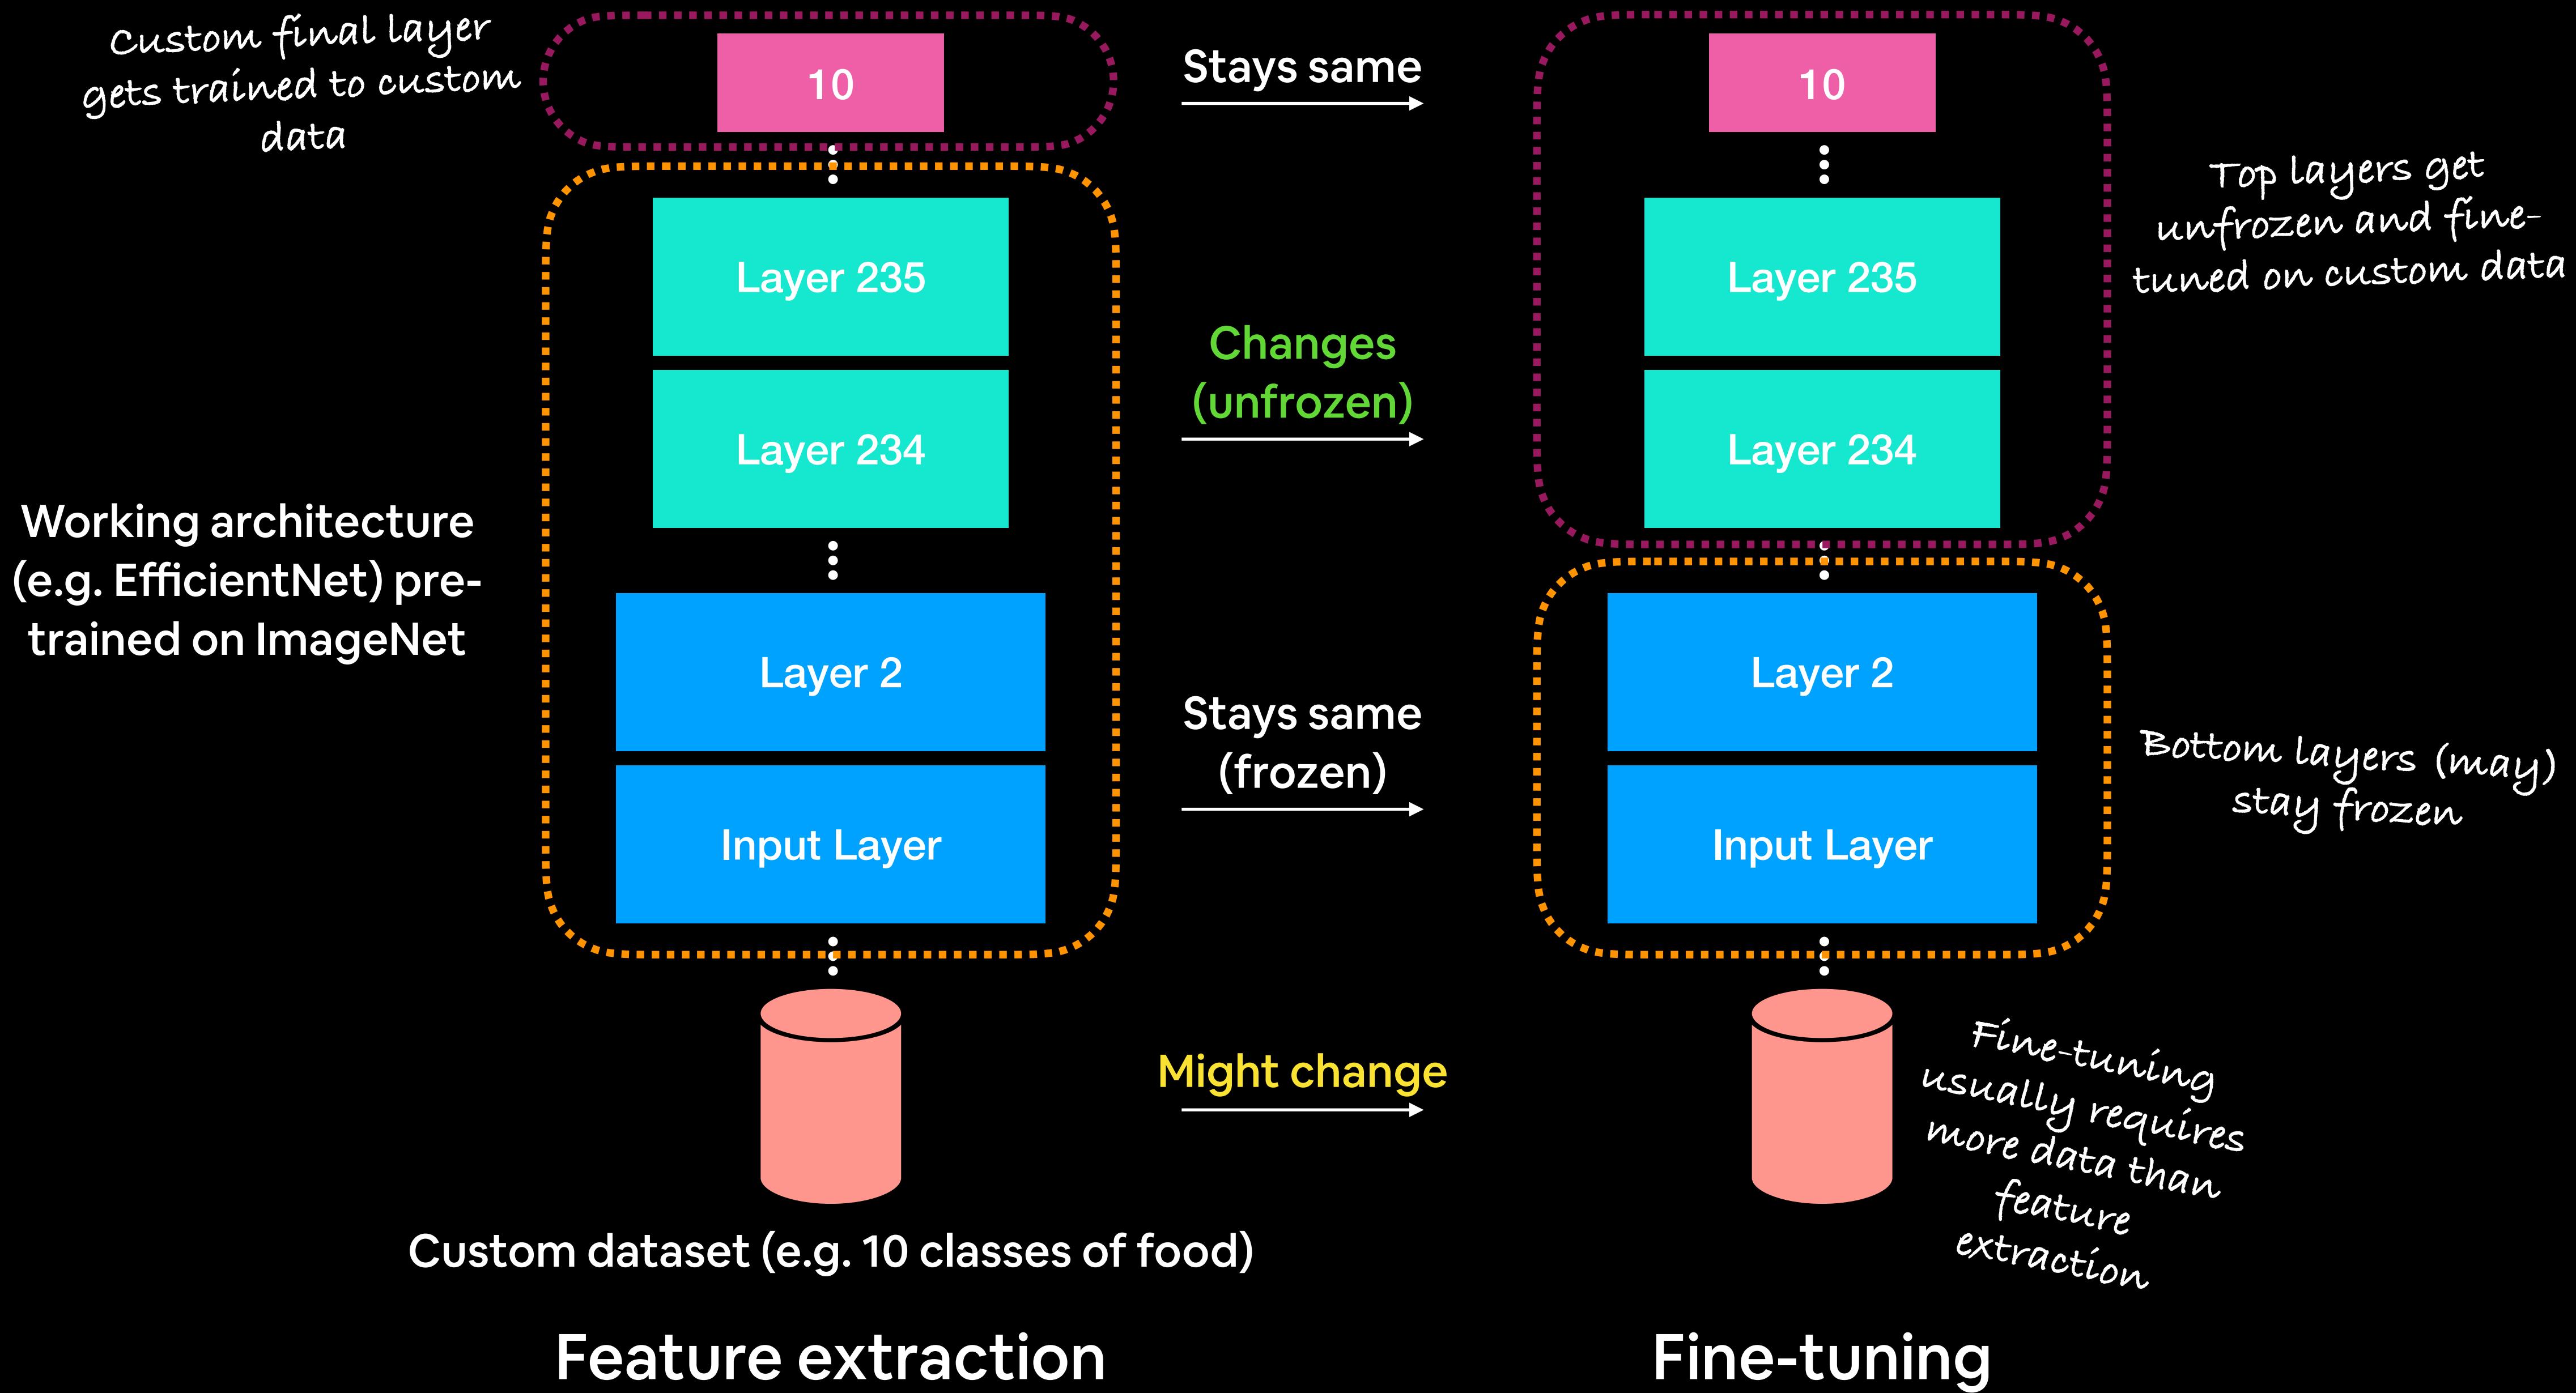
# Download base model and freeze underlying layers
base_model = tf.keras.applications.EfficientNetB0(include_top=False)
base_model.trainable = False

# Create functional model
inputs = layers.Input(shape=INPUT_SHAPE)
x = base_model(inputs, training=False)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dense(10)(x)
# Mixed precision requires output layer to have dtype=tf.float32 3
outputs = layers.Activation("softmax", dtype=tf.float32)(x)
model = tf.keras.Model(inputs, outputs)

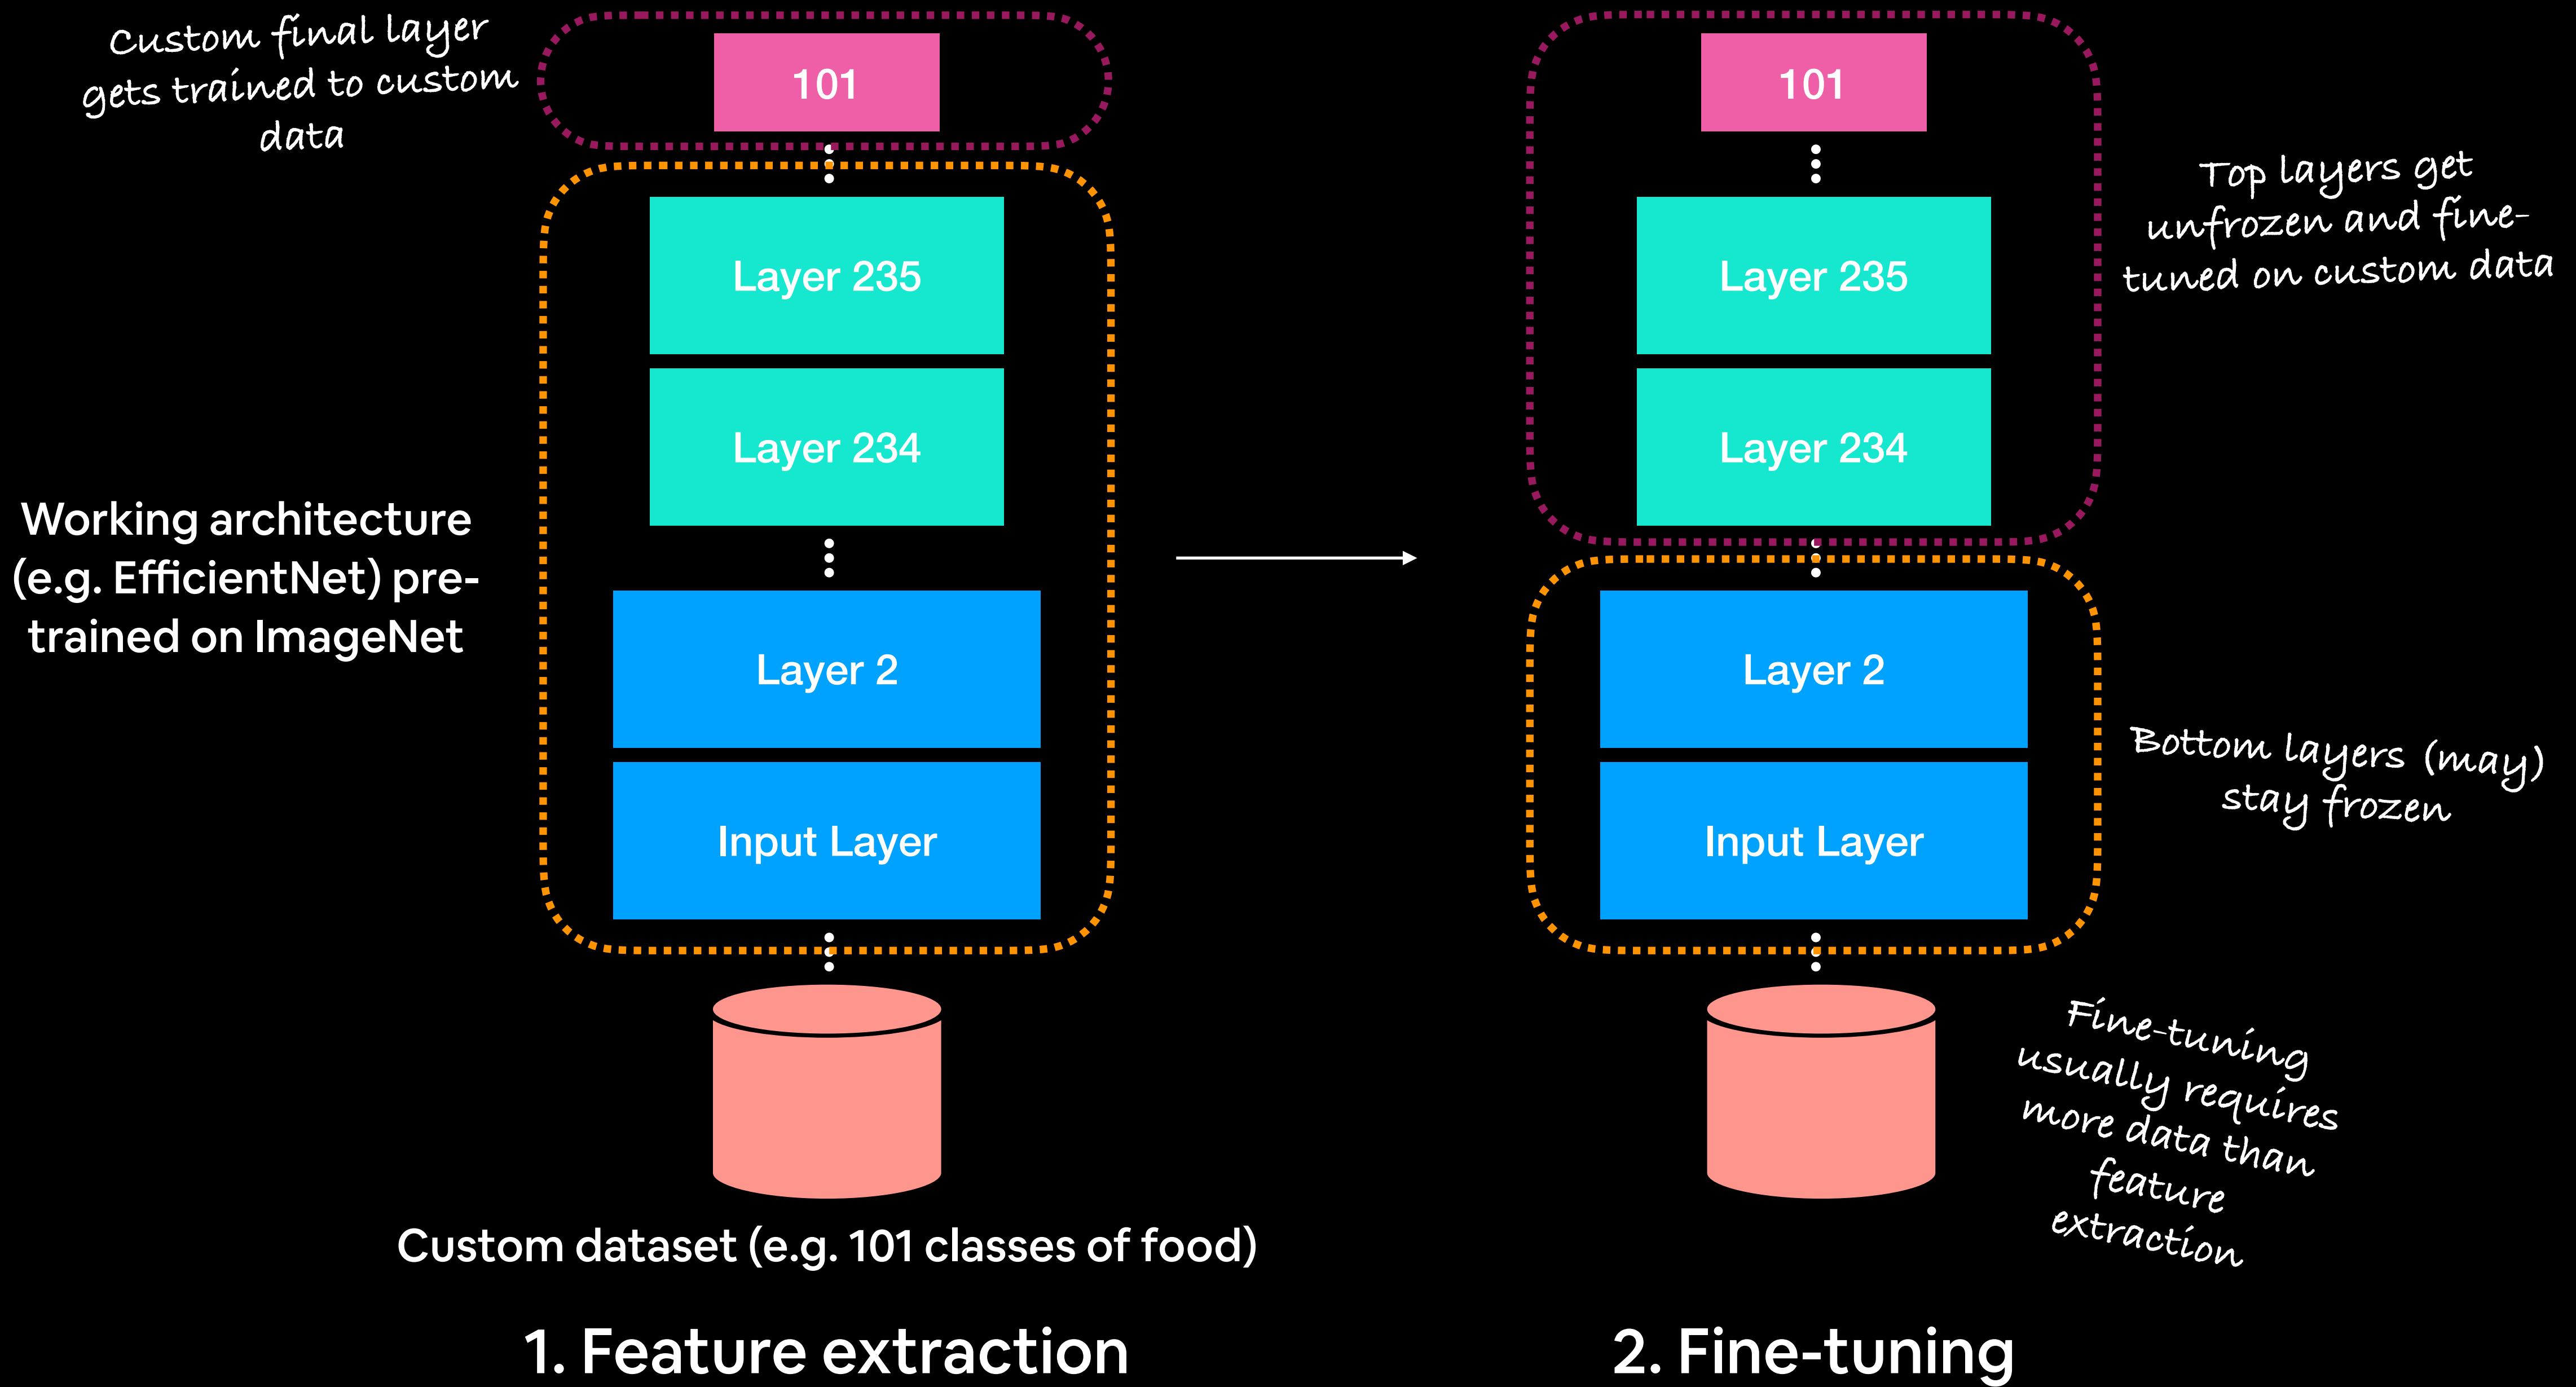
# Compile
model.compile(loss="sparse_categorical_crossentropy",
              optimizer=tf.keras.optimizers.Adam(),
              metrics=["accuracy"])
```

With mixed precision

Feature extraction vs. Fine-tuning



Feature extraction -> Fine-tuning



Performance of fast data fetching vs non-data fetching

	Prefetch and mixed precision	No prefetch and no mixed precision
Time per epoch	~280-300s	~1127-1397s

Results from fine-tuning  Food Vision Big™ on Food101 dataset using an EfficientNetB0 backbone using a Google Colab Tesla T4 GPU.

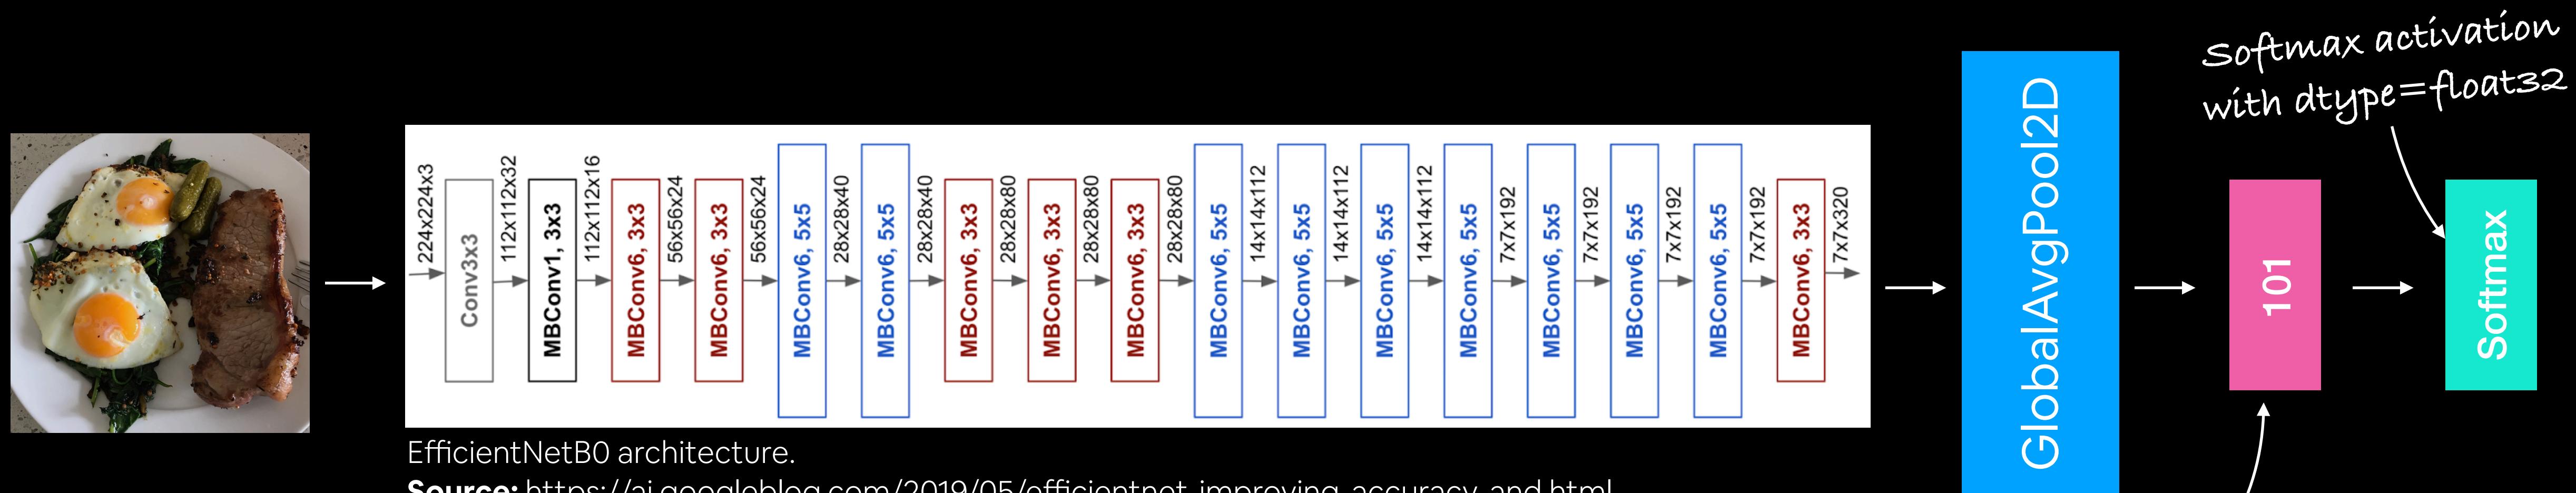
Model we've created,

(for mixed precision
training)

Input data

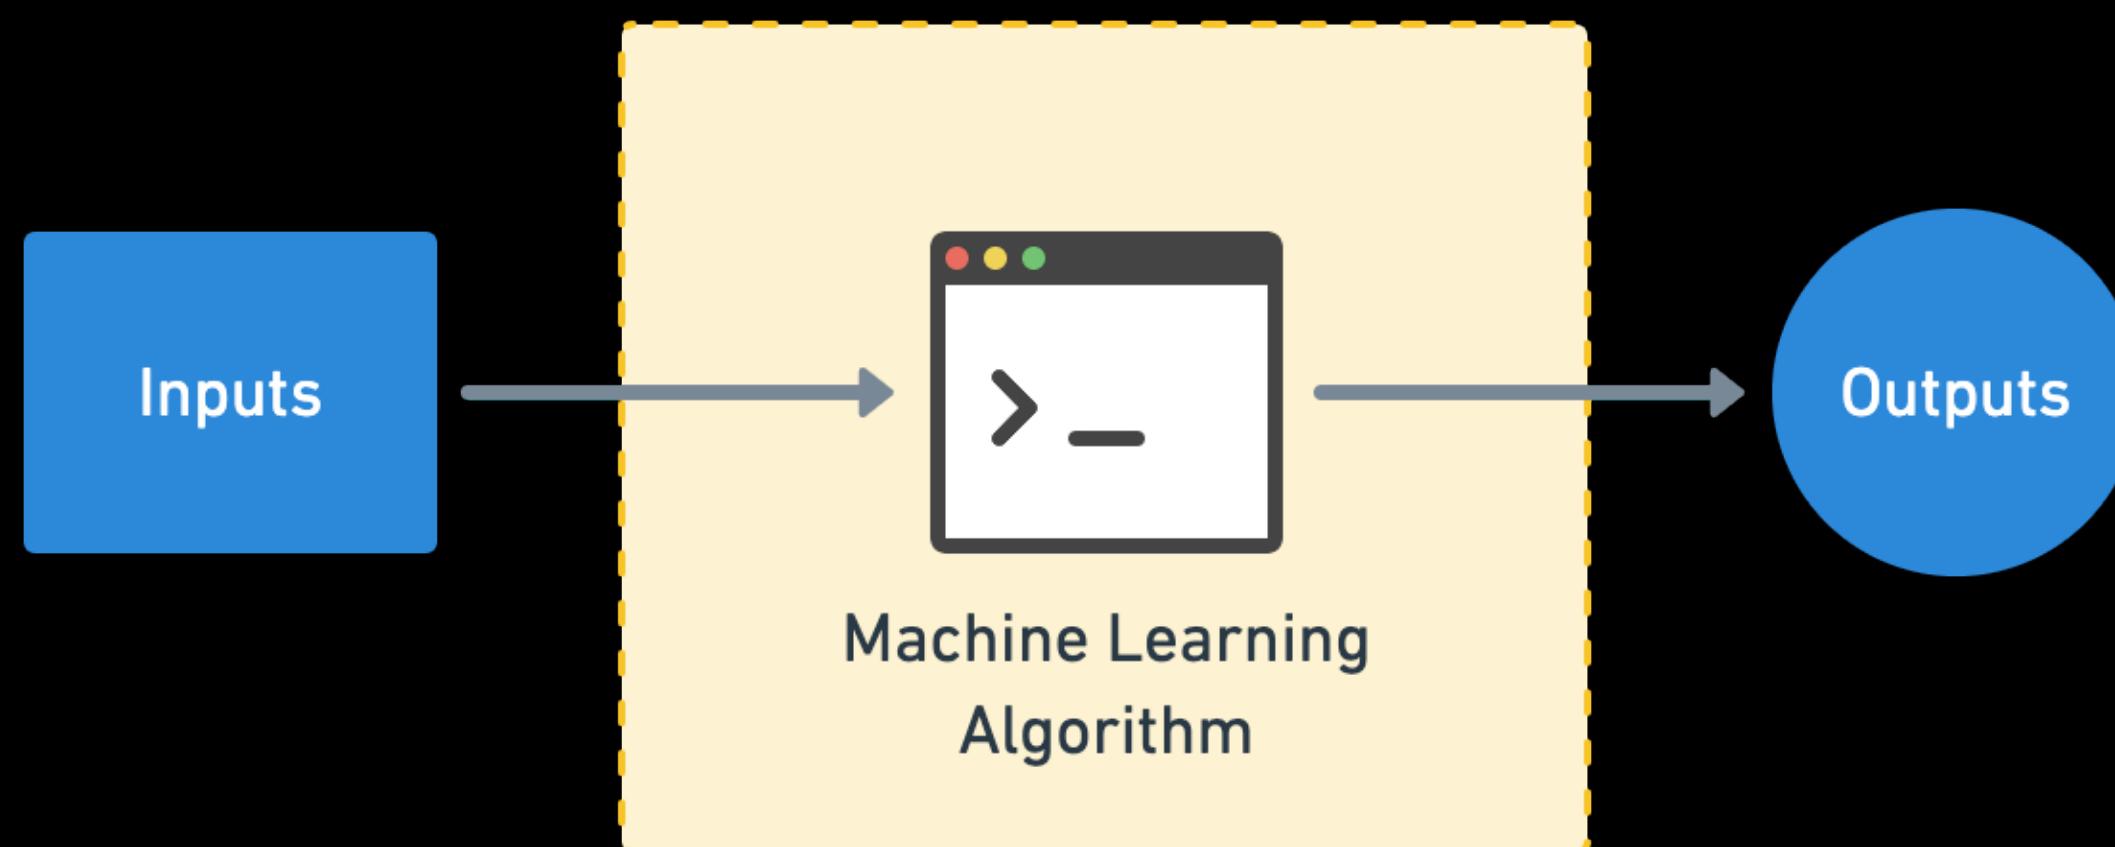
Model

Output



EfficientNetB0 architecture.

Source: <https://ai.googleblog.com/2019/05/efficientnet-improving-accuracy-and.html>



(some common)

Classification evaluation methods

Key: **tp** = True Positive, **tn** = True Negative, **fp** = False Positive, **fn** = False Negative

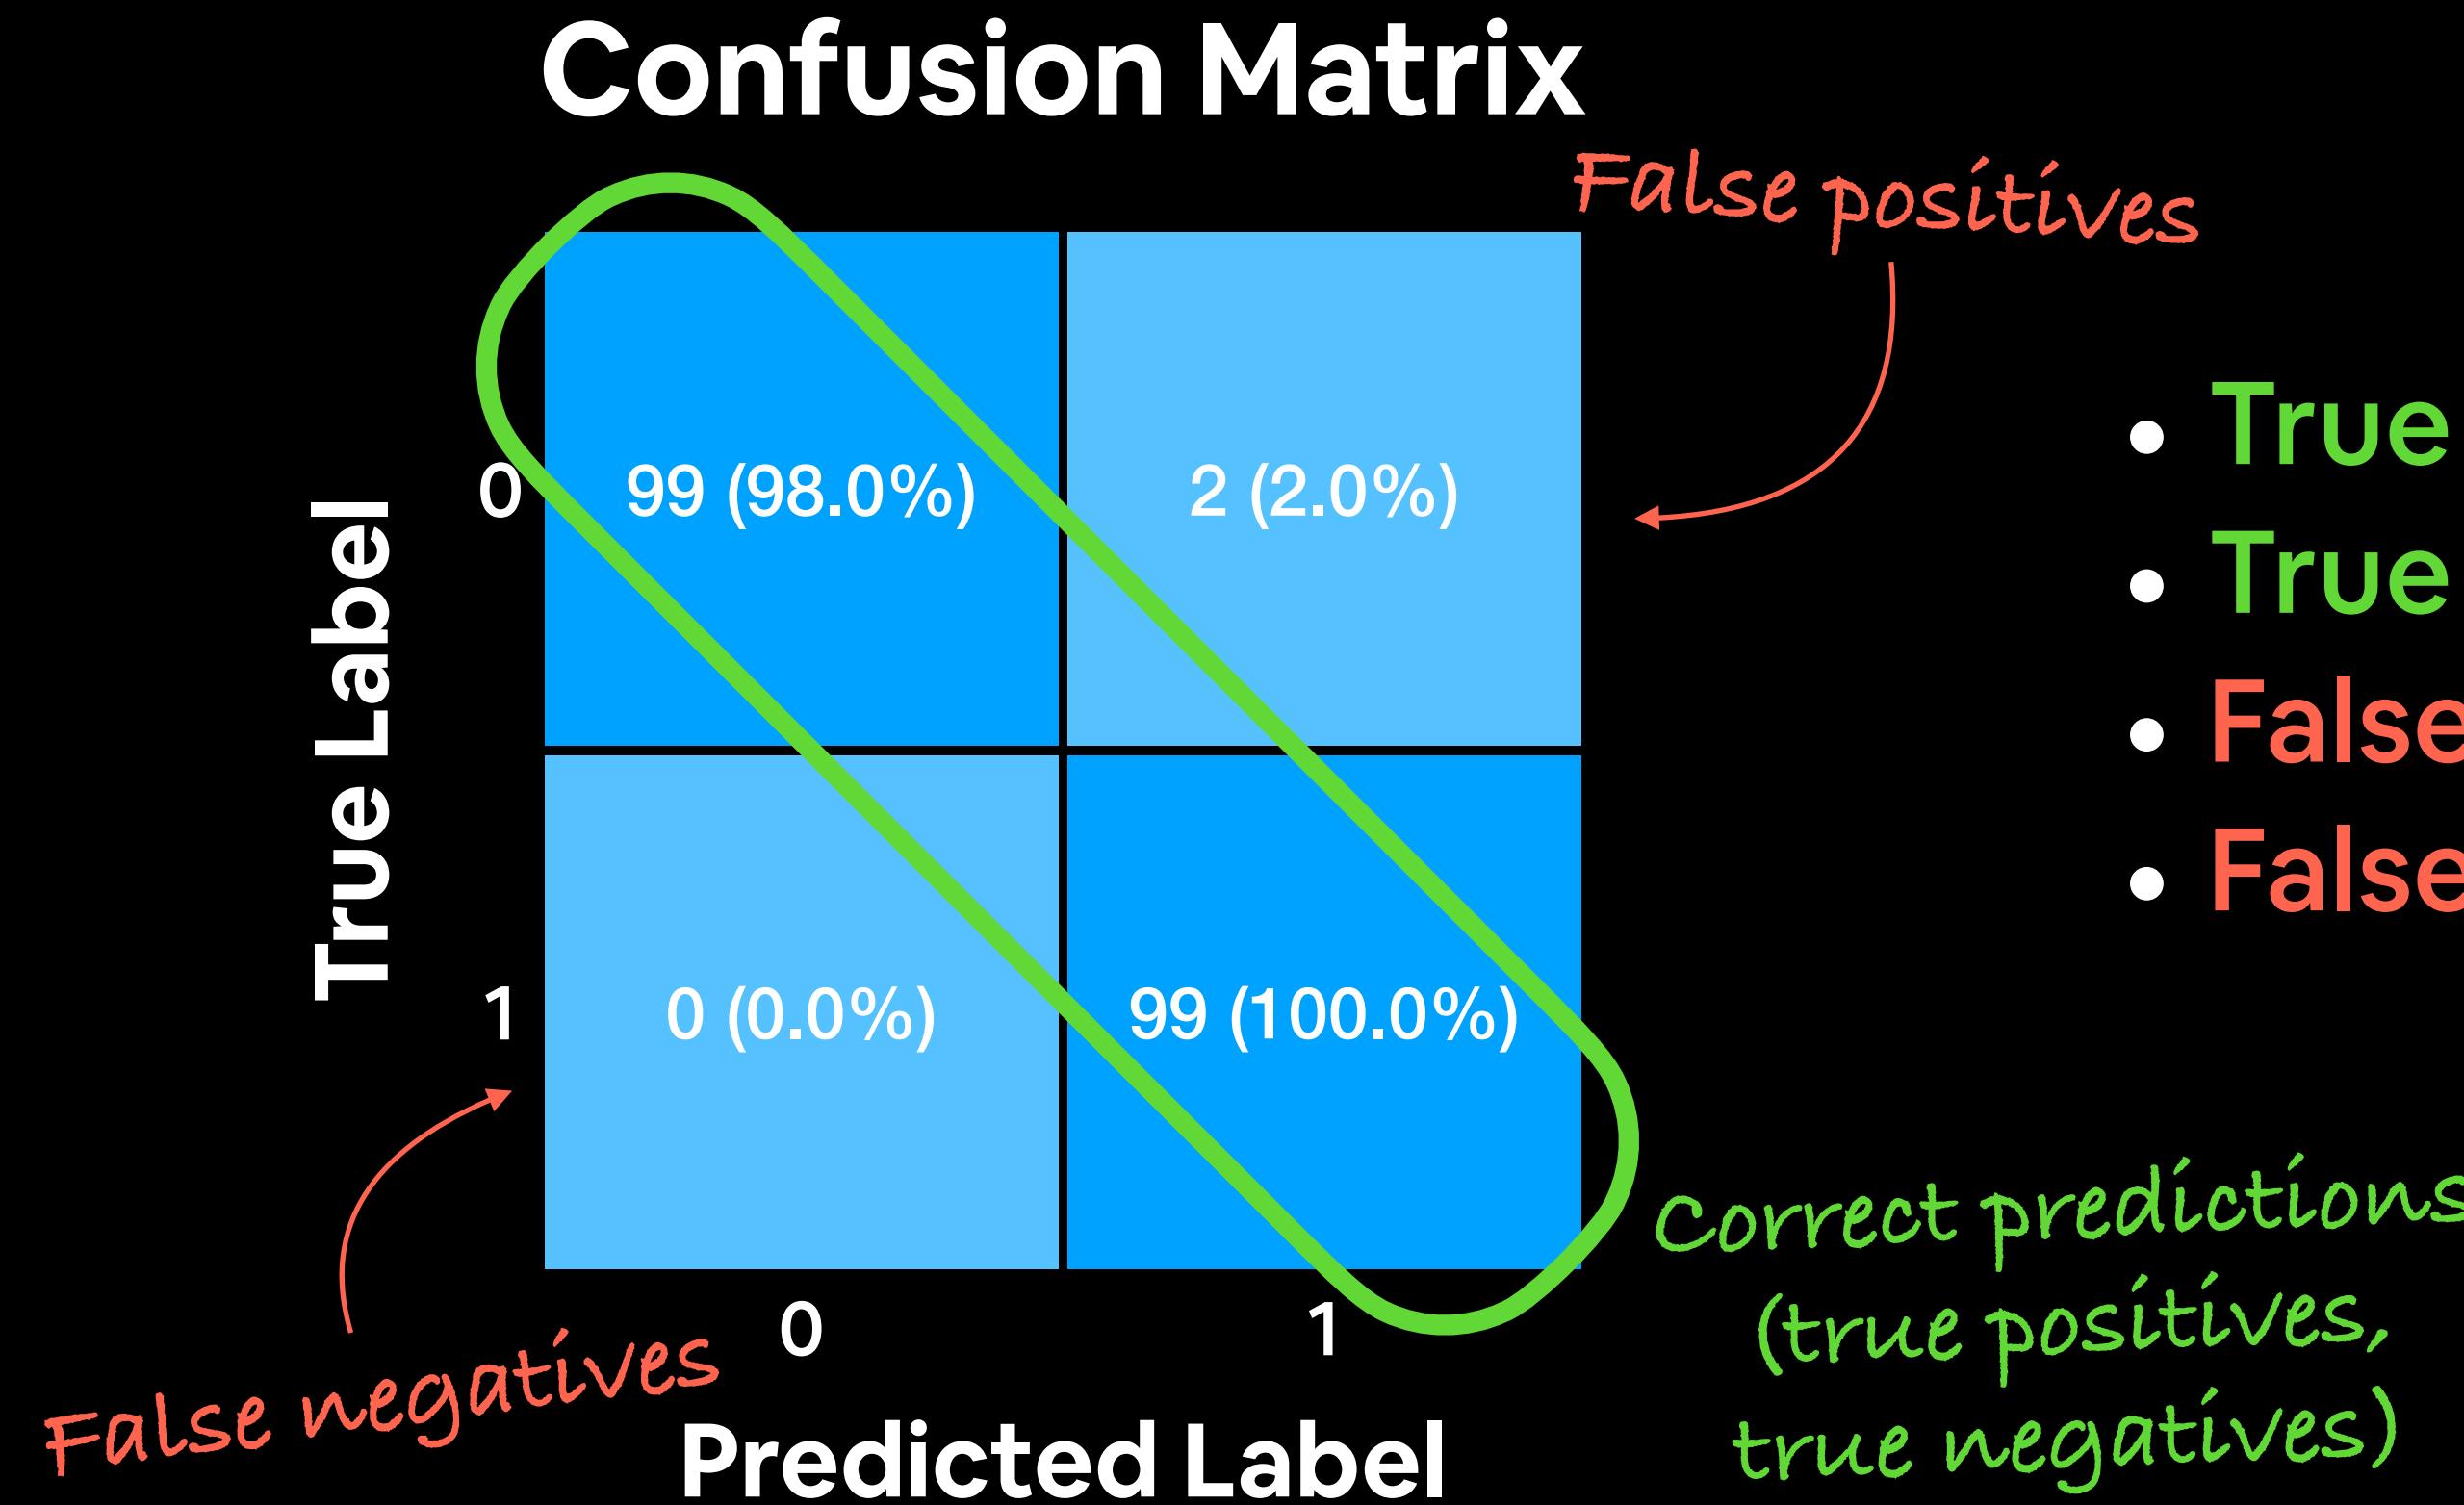
Metric Name	Metric Forumla	Code	When to use
Accuracy	Accuracy = $\frac{tp + tn}{tp + tn + fp + fn}$	<code>tf.keras.metrics.Accuracy()</code> or <code>sklearn.metrics.accuracy_score()</code>	Default metric for classification problems. Not the best for imbalanced classes.
Precision	Precision = $\frac{tp}{tp + fp}$	<code>tf.keras.metrics.Precision()</code> or <code>sklearn.metrics.precision_score()</code>	Higher precision leads to less false positives.
Recall	Recall = $\frac{tp}{tp + fn}$	<code>tf.keras.metrics.Recall()</code> or <code>sklearn.metrics.recall_score()</code>	Higher recall leads to less false negatives.
F1-score	F1-score = $2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$	<code>sklearn.metrics.f1_score()</code>	Combination of precision and recall, usually a good overall metric for a classification model.
Confusion matrix	NA	Custom function or <code>sklearn.metrics.confusion_matrix()</code>	When comparing predictions to truth labels to see where model gets confused. Can be hard to use with large numbers of classes.

Finding the most wrong predictions

- A good way to inspect your model's performance is to **view the wrong predictions with the highest prediction probability (or highest loss)**
- Can reveal insights such as:
 - Data issues (wrong labels, e.g. model is right, label is wrong)
 - Confusing classes (get better/more diverse data)



Anatomy of a confusion matrix



- **True positive** = model predicts 1 when truth is 1
- **True negative** = model predicts 0 when truth is 0
- **False positive** = model predicts 1 when truth is 0
- **False negative** = model predicts 0 when truth is 1